



Managing Machine Learning Project with Visual Studio

Dr. Ridi Ferdiana
Microsoft MVP
Microsoft Certified Trainers
Microsoft Certified for DevOps Expert



<https://mic.ms/ridi>




I am not an expert on Machine Learning, so I can say “I don’t know”

– Ridi

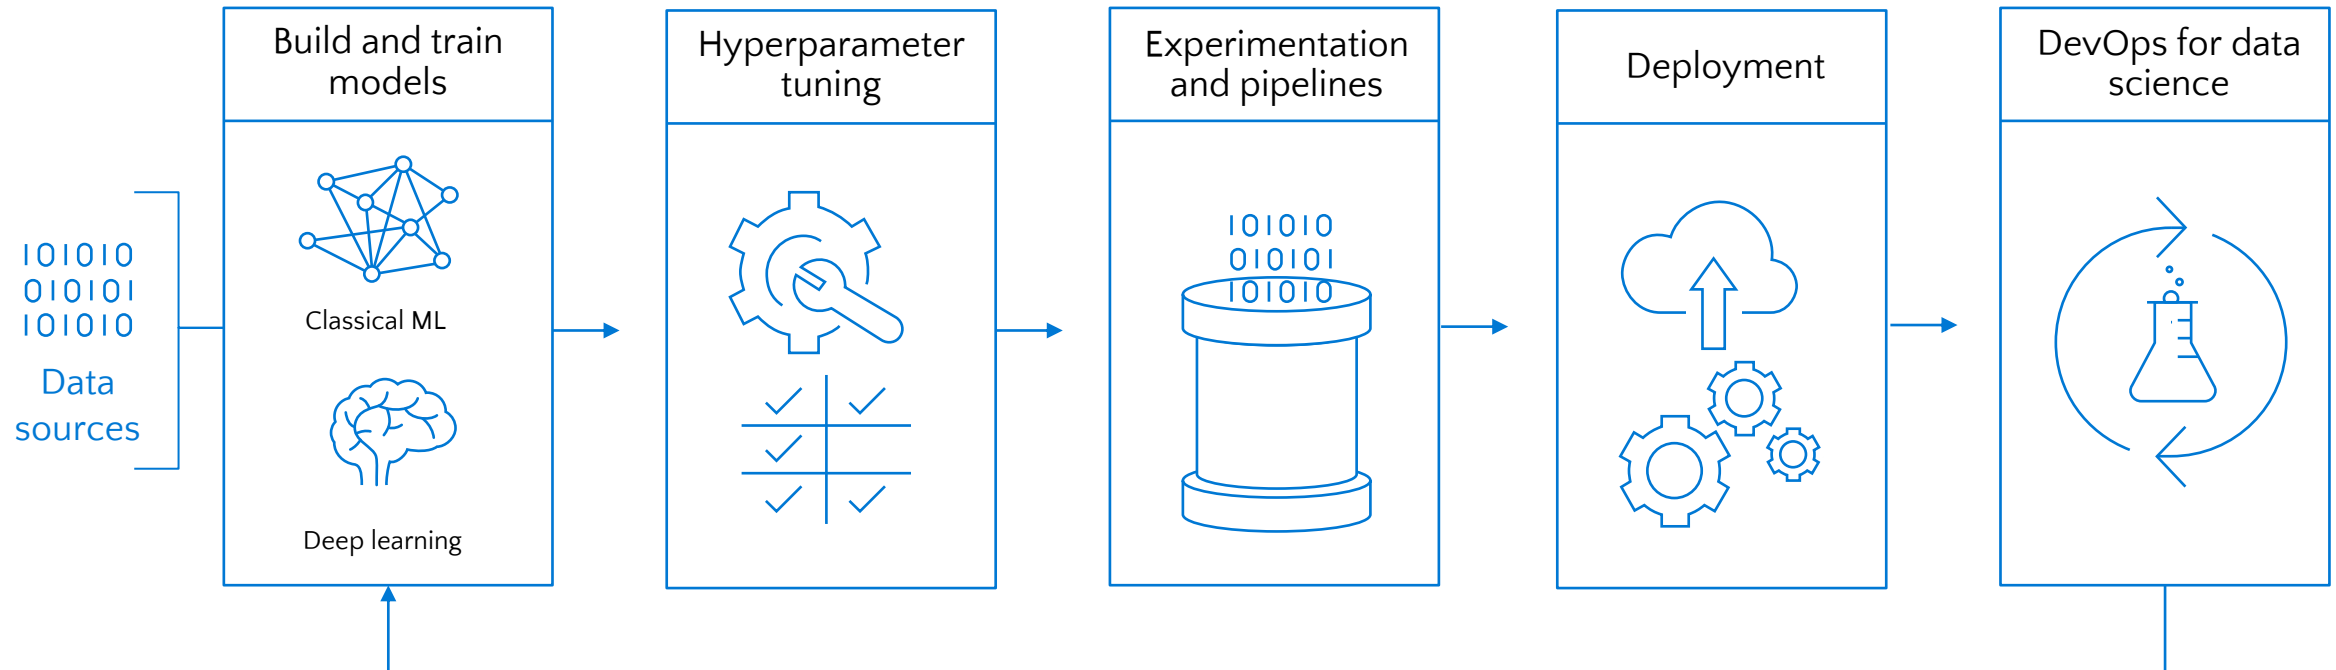
Our Agenda in 30 Minutes

- Data Science Project on Microsoft Ecosystem
- DevOps for Machine Learning (MLOps)

Data Science Project on Microsoft Ecosystem




Building blocks for a Data Science Project



Company point of view



Serving business users and end users with **intelligent** and **dynamic** applications



Build a unified and usable
data pipeline



Train ML and DL models to
derive insights



Operationalize models and
distribute insights at scale

Machine Learning on Microsoft Ecosystem

Domain Specific Pretrained Models

To reduce time to market



Vision



Speech



Language



Search

Familiar Data Science Tools

To simplify model development



PyCharm



Jupyter



Visual Studio Code



Command line

Popular Frameworks

To build machine learning and deep learning solutions



PyTorch



TensorFlow



Scikit-Learn



ONNX

Productive Services

To empower data science and development teams



Azure Databricks



Azure Machine Learning



Machine Learning VMs

Powerful Hardware

To accelerate deep learning



CPU



GPU



FPGA









From the Intelligent Cloud to the Intelligent Edge



Cognitive Services






Vision

Recognize, identify, caption, index, and moderate your pictures, videos, and digital ink content.

-  [Computer Vision](#)
-  [Custom Vision](#)
-  [Face](#)
-  [Form Recognizer](#)
-  [Ink Recognizer \(Retiring\)](#)
-  [Video Indexer](#)




Language

Allow your apps to process natural language with pre-built scripts, evaluate sentiment and learn how to recognize what users want.

-  [Immersive Reader](#)
-  [Language Understanding \(LUIS\)](#)
-  [QnA Maker](#)
-  [Text Analytics](#)
-  [Translator](#)





Speech

Convert speech into text and text into natural-sounding speech. Translate from one language to another and enable speaker verification and recognition.

-  [Speech service](#)
-  [- Customize with Speech Studio](#) 


Decision

Build apps that surface recommendations for informed and efficient decision-making.

-  [Anomaly Detector](#)
-  [Content Moderator](#)
-  [Metrics Advisor \(preview\)](#)
-  [Personalizer](#)


Search

Add Bing Search APIs to your apps and harness the ability to comb billions of webpages, images, videos, and news with a single API call.

-  [Bing Search APIs hub page](#)




Cognitive Service Containers

Container support in Azure Cognitive Services allows developers to use the same rich APIs that are available in Azure, and enables flexibility in where to deploy and host the services that come with Docker containers.

-  [Cognitive Services containers documentation](#)
-  [What are Cognitive Service containers?](#)
-  [Create containers for reuse](#)
-  [Deploy and run container on Azure Container Instance](#)
-  [Deploy to Azure Kubernetes Service](#)
-  [Use Docker Compose to deploy multiple containers](#)
-  [Cognitive Services containers FAQ](#)

Cognitive Services for Big Data

The Cognitive Services for Big Data allows developers to use cognitive services within databases to apply intelligent algorithms over large datasets

-  [Cognitive Services for Big Data](#)
-  [Use Cognitive Services within Azure Databricks and Synapse Analytics](#)
-  [Python Samples](#)

If you really need a custom solution

Empower data science and development teams



Integrated data science & data engineering teams

Desktop solutions not adequate

Need a unified big data & machine learning solution



Azure Databricks

+



Azure Machine Learning
service



Individual data scientists

Desktop solutions adequate

Need cloud for sporadic compute needs

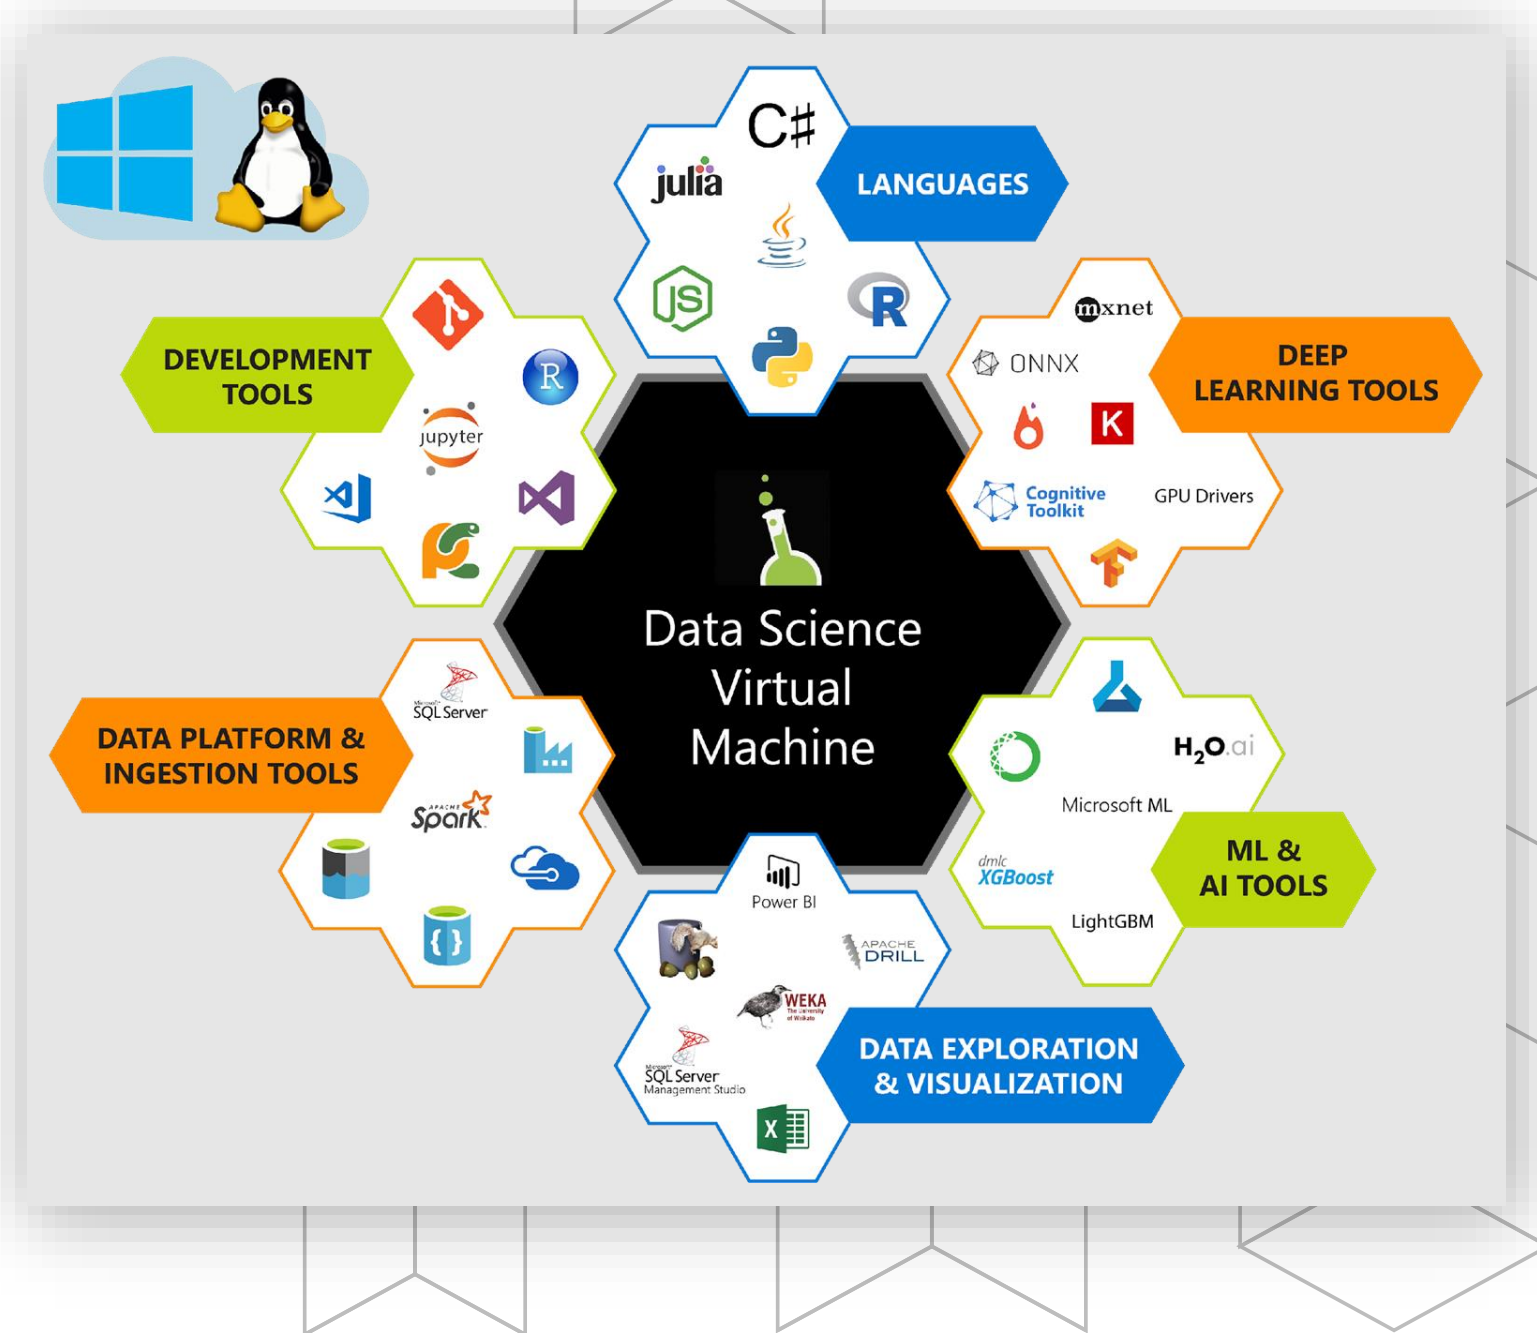


Machine Learning VMs

Data Science Virtual Machines (DSVM)

Pre-Configured environments in the cloud for
Data Science & AI Modeling,
Development & Deployment.


Samples to get started




Why Data Science VMs?




Elastic analytics
desktop in the
cloud



Get started quickly
on Azure Machine
Learning



Pre-setup
environment at
work and school



Enterprise Ready
and globally
available

DEMO

Preparing Data Science Virtual Machine on Azure



Azure Machine Learning Service

Set of Azure Cloud
Services



Python
SDK

That enables you to:


- ✓ Prepare Data
- ✓ Build Models
- ✓ Train Models

- ✓ Manage Models
- ✓ Track Experiments
- ✓ Deploy Models



Azure Databricks

Fast, easy, and collaborative Apache Spark™-based analytics platform



Increase productivity



Build on a secure, trusted cloud



Scale without limits



Built with your needs in mind

- Optimized Apache Spark environment
- Collaborative workspace
- Integration with Azure data services
- Autoscale and autoterminate
- Optimized for distributed processing
- Support for multiple languages and libraries

Seamlessly integrated with the Azure Portfolio



Azure Machine Learning service

Bring AI to everyone with an end-to-end, scalable, trusted platform



Boost your data science productivity



Increase your rate of experimentation



Deploy and manage your models everywhere



Built with your needs in mind

- Automated machine learning
- Managed compute
- DevOps for machine learning
- Simple deployment
- Tool agnostic Python SDK
- Support for open source frameworks

Seamlessly integrated with the Azure Portfolio

Powerful infrastructure

Accelerate deep learning



CPUs

General purpose machine learning


D, F, L, M, H Series



GPUs

Deep learning

N Series



FPGAs

Specialized hardware accelerated deep learning

Project Brainwave

Optimized for flexibility


Optimized for performance

DEMO

Azure Machine Learning Studio




DevOps for Machine Learning



DevOps loop for data science

Prepare



Prepare
Data



Experiment



...



Build model
(your favorite IDE)



Train &
Test Model



Register and
Manage Model

Deploy

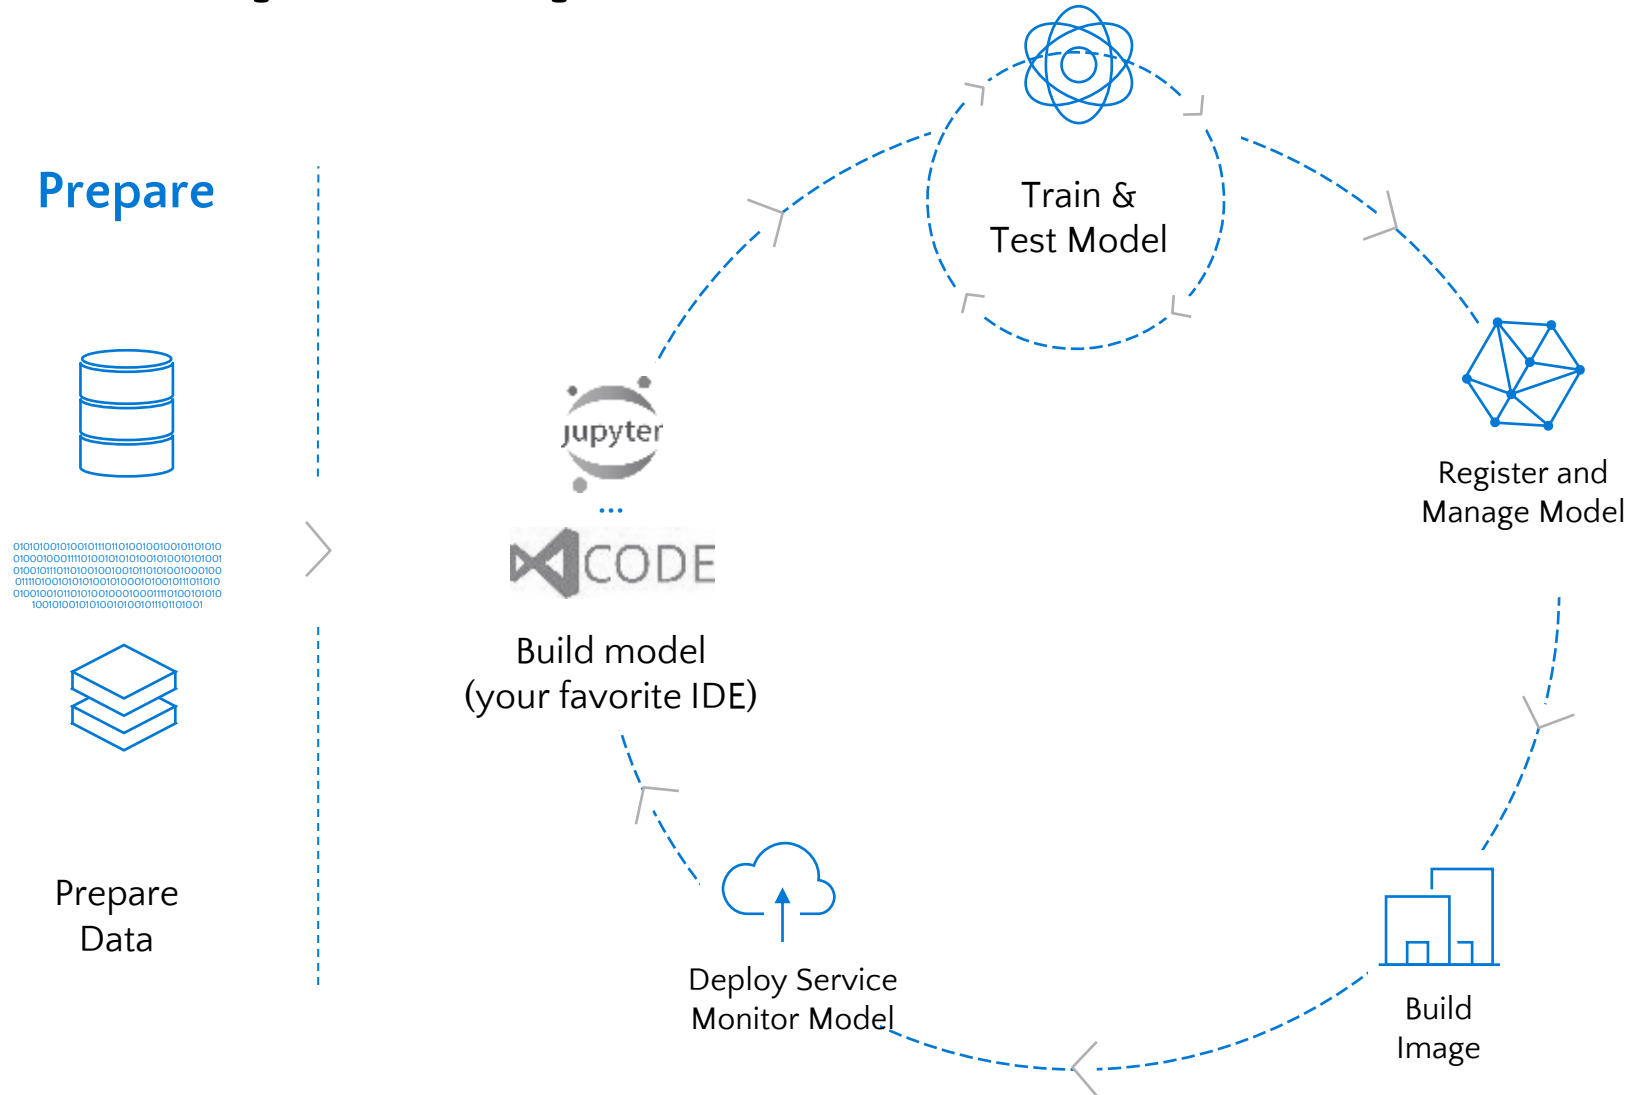


Build
Image



Deploy Service
Monitor Model

DevOps loop for data science




Advantages of Azure ML Pipelines




Unattended runs

Schedule a few steps to run in parallel or in sequence to focus on other tasks while your pipeline runs




Tracking and versioning

Name and version your data sources, inputs and outputs with the pipelines SDK



Reusability

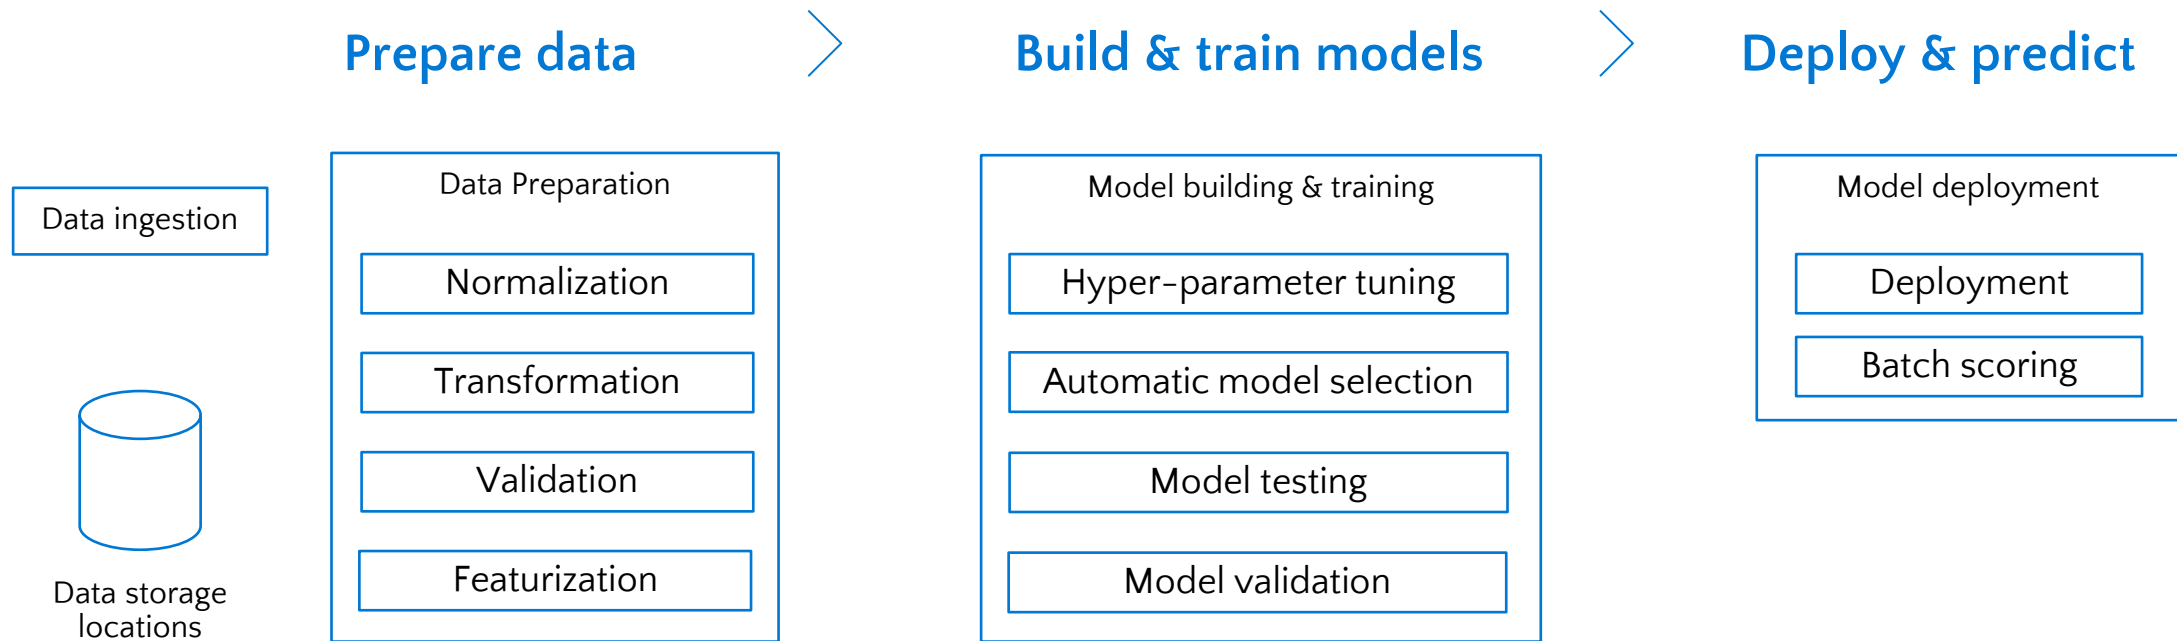
Create templates of pipelines for specific scenarios such as retraining and batch scoring



Mixed and diverse compute

Use multiple pipelines that are reliably coordinated across heterogeneous and scalable computes and storages


Azure Machine Learning Pipelines



Deployments to Compute Targets

| Compute target | Deployment type | Description |
|--|-----------------|--|
| Azure Container Instances (ACI) | Web service | Fast deployment. Good for development or testing. |
| Azure Kubernetes Service (AKS) | Web service | Good for high-scale production deployments. Provides autoscaling, and fast response times. |
| Azure IoT Edge | IoT module | Deploy models on IoT devices. Inferencing happens on the device. |
| Field-programmable gate array (FPGA) | Web service | Ultra-low latency for real-time inferencing. |

Case Study





Requirement Azure ML Studio

- Azure Data Blob
- Compute Instance (Standard_DS1_v2) for Notebook
- Compute Cluster (Standard_DS11_v2) for Pipeline

Collecting the Dataset

Dataset Collection

- **Data Resource:** Youtube Audio collected by Audioset.
- **Data Type:** An audio file (.flac) with a duration of 10 seconds
- **Amount of data :** 484
- **Data Label:** happy, sad, relaxed, tense, angry.

Dataset Transformation

- **Transformation Method:** Mel-Frequency Cepstral Coefficients (MFCC)
- **Data training and testing:** Divide by the ratio of 8: 2
- **Data stored in Json form**

CatSign_Algorithm > Datasets > Catsign

Catsign

Version 1 (latest) ▾

Details

Consume

Explore

Models

Refresh ▶ Generate profile ★ Unregister 📄 New version ▾

Preview

Profile

Number of columns: 2 Number of rows: 50 (of 484)

| Id | feature | class_label |
|----|---------------------------------------|-------------|
| 1 | [-498.1416625977, 131.9505310059, ... | Caterwaul |
| 2 | [-341.6528625488, 171.3036346436, ... | Caterwaul |
| 3 | [-435.2915039062, 191.8328552246, ... | Caterwaul |
| 4 | [-340.2807922363, 170.4213562012, ... | Caterwaul |
| 5 | [-361.5213928223, 104.8150253296, ... | Caterwaul |
| 6 | [-299.4932861328, 133.0233612061, ... | Caterwaul |
| 7 | [-229.6245269775, 84.0176544189, -... | Caterwaul |
| 8 | [-147.2395935059, 73.9755325317, -... | Caterwaul |
| 9 | [-539.1246337891, 118.9978408813, ... | Caterwaul |
| 10 | [-190.7236633301, 66.3537216187, -... | Caterwaul |

Importing dataset to AMLS

- Azure Machine Learning Studio datasets can be added via the menu Assets → Datasets.
- The dataset will be stored on Azure Blob Storage.
- As the data changes, you can add new versions to existing dataset assets.
- The registered dataset can be used in Automated ML and Machine Learning Designer

Make a CatSign model with ML Designer

01

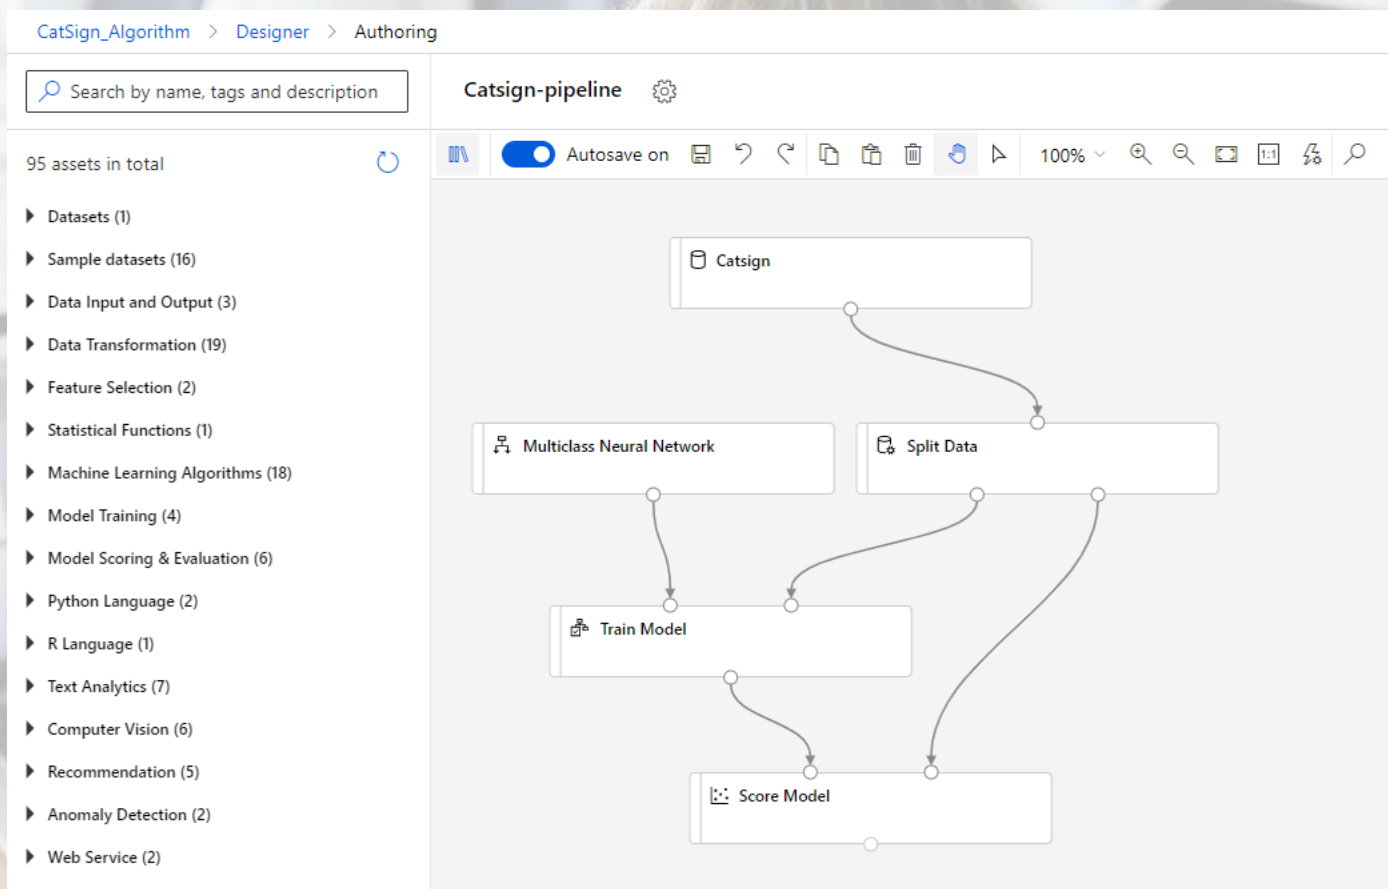
Start by creating a pipeline via the Pipelines menu on the sidebar.

02


Select the Compute Cluster created for the compute source

03

Select datasets, model algorithms, and tasks for training and evaluation



Summary






Azure Machine Learning service

Bring AI to everyone with an end-to-end, scalable, trusted platform



Boost your data science productivity



Increase your rate of experimentation



Deploy and manage your models everywhere



Built with your needs in mind

- Automated machine learning
- Managed compute
- Simple deployment
- DevOps for machine learning
- Support for open source frameworks
- Tool agnostic Python SDK

Seamlessly integrated with the Azure Portfolio

Resources beyond this AI Airlift

The background features several orange wireframe cubes of varying sizes and orientations, scattered across the right side and bottom of the slide.

Azure Machine Learning with Azure Databricks

<https://aka.ms/aml-notebook-databricks-e2e>

Azure Notebooks

<https://notebooks.azure.com/azureml/projects/azureml-getting-started>

Azure ML Docs

<https://docs.microsoft.com/en-us/azure/machine-learning/service/>